

# Sybase Masterclass: Update statistics like we mean business

# Contents

---

- About the author
- Why this presentation?
- Why do database statistics matter so much?
- Why are update stats difficult?
- Update stats... in the right places
- Update stats... at the right times
- Update stats... in the right ways

# About the author

---

- Sybase Australia 1996 – 2003; youngest ever Principal Consultant
- After working with Sybase tech for 20+ years... is no longer young
- Principal Database Engineer at Prima Donna Consulting
- Based in London, UK, and Melbourne, Australia
- International Sybase User Group Board of Directors since 2010
- UK Sybase User Group Board of Directors since 2019
- Not a lawyer – no charge for emails!
- Improves client bottom lines by ~USD\$10M/month every month

# Why this presentation

---

- Update stats are one of the most frequently run DBA housekeeping
  - ... and yet I have never seen a site with no stats problems ever
- It looks easy and simple
  - It is a dangerously misleading topic
- Consequences for getting it wrong: high-priority incidents
- This presentation is practical and pragmatic
  - The theory and description of the stats are in the manuals

# Why do database statistics matter so much?


---

- It's all about the ASE optimiser
- Unless a statement is one table with no indexes, decisions must be made
- Which access method? Table scan? Index? Which one?
- Which join order?
- Which join method?
- Even a statement with one table with no indexes still needs decisions
  - Parallel query? Lock type? Lock promotions? Fetch-and-discard?

# Why do database statistics matter so much?

- These decisions crucially rely on knowing what the data looks like:

- number of pages
- number of rows
- density – how unique?
- distribution – how even?



# Why do database statistics matter so much?

---

- Number of pages
  - Access methods: how bad is a table scan?
  - Access methods: how good is this index?
- Number of rows
  - Join types: nested loop join starts getting expensive
  - Join order: prefer the smallest table innermost for any loops

# Why do database statistics matter so much?

---

- Density
  - Uniqueness: how many rows will match this value (or range)?
- Distribution
  - How evenly distributed are the values in this column?
  - If we don't know then our predictions get much worse



# The cost of missing or incorrect stats

---

- Missing stats = magic numbers
  - = 10%
  - <, > 25%
  - <=, >= 33%
- Incorrect stats = ASE is misinformed and makes wrong decisions
- Q: Are these scenarios really that bad?
- A: OMG yes
  - One incorrect table scan or join order can bring ASE to its knees

# Why this is so difficult to get right

---

- OK, so missing or wrong stats are bad: run them all the time!
- Update stats consume time & resources, and impact database use
  - There are clear downsides to running them too often
- i.e. this is a Goldilocks problem
  - Run update stats at least as often as required to prevent problems...
  - ... and no more often than that

# The cost of update stats too much / too often

---

- Exact cost of update stats depends on which kind of update, but...
  - I/O                      table; index(es)
  - memory                procedure cache; data cache
  - tempdb                named tempdbs are bound to login or application  
(how many of you bind sa/DBA to small tempdb?)
  - CPU                    small but detectable
  - locks                  number of locks; type of locks

# Update stats... in the right places

---

- Three obvious and yet profound truths:
  - Statistics are only necessary when they will change a query plan
  - Statistics don't change if the data doesn't change
  - This means we can't get update stats decisions right without:
    - Knowing all queries against the data, and
    - Knowing what data changes, how, and when

# Update stats... in the right places

---

- Statistics are only necessary when they will change a query plan
  - If there are no indexes on a table, there are no index decisions
  - If there is only one table in a join, there are no join decisions
  - If parallel is disabled (or ruled out), there are no parallel decisions
- Put another way, any query that always table scans a single table...
  - ... doesn't need update stats

# Update stats... in the right places

---

- Statistics don't change if the data doesn't change
  - Read-only tables need stats updated only once
  - Seldom-written tables need stats updated only seldom
  - Append-only tables need update stats only on what is appended
    - Partitions, and only update stats on changed partition(s)?
    - Requires local indexes = problems with global uniqueness

# Partitions and local indexes

---

- Semantic partitions = premium license
- Allows us to define which partition any row lives in
- Allows us to index per partition (local) as well of instead of global
- uniqueness is difficult; needs either:
  - At least one unique global index (defeats the purpose?)
    - But no global index makes quality of global stats worse
  - Partition key must appear in local unique indexes
- If you can make this work, only update stats on changed partitions(s)

# Partitions and local indexes

---

- Fake it with UNION ALL in VIEW?
- Main table + spill table + infrequently batch job to move spill to main
- Most of the time only run update stats on spill table
- Only ever update stats on main table when spill is moved to main
  - Monthly? Quarterly?
  - We don't care if this takes a long time (in terms of locking)
  - We only ever read from main; all writes go to spill table
- Query performance “should” be close; test it first



# Update stats... in the right places

---

- The elephant in the room!
- We don't need to **ever** run update stats in PROD
- Most sites do nightly refresh PROD to UAT / PERF TEST / *somewhere*
  - And if you don't, it's time to start
- Run update stats on one-day-old PROD copy
- optdiag out from UAT
- optdiag in to PROD

# Update stats... in the right times

---

- When to run? As much as is needed, and no more than that
- So when is it needed?
  - When data is changing in ways that change plans
  - When stats are stale
  - When stats are missing
  - When stats weren't updated even though we updated them
- Update stats for free...

# When data is changing: datachange()

---

- Counts inserts and deletes
  - Update = delete + insert = 2 “changes”
  - UPDATE tableA SET colX = colX = 100% “change”
- Expressed as a percentage (can be > 100%)
  - Interpret in context of rowcount
  - 1% x 100 rows = who cares; 1% x 10B rows = we care a lot

# When data is changing: datachange()

---

- Not transactional, rollbacks don't reverse the increment
  - Can be lost from memory
- Refreshed during any polite shutdown
- Refreshed regularly based on “sysstatistics flush interval”...
  - Which defaults to 0 !!!
  - Units are minutes
  - Set to something sensible like 15 or 60
  - You definitely don't need this set to 1

# When stats are stale

---

- Probably the most common cause of stats-related issues
  - Data changes can affect optimiser decisions in non-linear ways
    - i.e. +1% data  $\Rightarrow$  200x slower because of sudden table scan
    - A clear sign at least some stats aren't updated often enough
- How would we tell?

# When stats are stale

---

- optdiag
- sp\_showoptstats
  - Based on sp\_\_optdiag but not as good (thanks Kevin, Raymond)
- select moddate from sysstatistics
- Remember stats are stored per column, not per index
- Check all the columns you care about
- How stale is too stale?
  - Often we can only find out the hard way...

# When stats are missing

---

- **DON'T RUN IN PRODUCTION!!!**
- set option `show_missing_stats`
- `sp_configure` “capture missing stats”
  - Does not display them anywhere; sends them to sysstatistics
  - **These are never deleted unless a DBA manually deletes them!**
  - Not human readable, but I've got your back, see next slides
- Only noticed when a query plan is generated
  - **Not captured for plans already in procedure / statement cache!**
- `#temp_tables` only reliably shown using set option `show_missing_stats`

# When stats are missing (out of the manuals)

```
create function sp_colidarray
    @colidarray varbinary(100)
    , @id int
returns varchar (1500)
as
declare @s varchar (1500)
    , @len int
    , @colid int
    , @colname longsysname
set @len = datalength(@colidarray)
while @len > 0
begin
    set @colid = convert(tinyint,
                        substring(@colidarray, @len, 1))

    set @colname = col_name(@id, @colid)
    set @s = @colname +
        case @s when NULL then NULL
              else ','
        end + @s
    set @len = @len - 2
end
return @s
```

```
create proc sp_decode_missing_stats
    @tabname varchar (100) = '%'
as
select Dbname = db_name()
    , Tabname = object_name(s.id)
    , NrRows = row_count(db_id(), s.id)
    , ColumnList =
        dbo.sp_colidarray(colidarray, s.id)
    , Captured = moddate
    , Occurs = convert(smallint, c0)
into #missing
from sysstatistics s
    , sysobjects o
where s.id = o.id
and object_name(s.id) like @tabname
and formatid = 110
and datalength(colidarray) > 0
exec sp_autoformat
    #missing,
    @orderby="order by 2"
```





# When stats are missing (out of the manuals)

```
exec my_db..sp_decode_missing_stats
```

Dbname	Tabname	NrRows	ColumnList	Captured	Occurs
-----	-----	-----	-----	-----	-----
mydb	my_table	7800	col_a	May 26 2009 8:26AM	14
mydb	your_table	5	col_b, col_c	May 26 2009 8:23AM	1
mydb	other_table	54998666	col_d	May 19 2009 10:03PM	202

# When stats are missing

---

- Don't run wild with missing stats
- Need to know how the tables and columns are used
- Need to prove that missing stats are generating bad plans
- We might still want stats even on unindexed columns
  - This can be cheaper than adding/modifying an index
  - More on this later
- You update them, you maintain them
  - Updating special stats only once may be worse than no stats

# When stats are missing

---

- Special case: systabstats
- Table-level stats such as page counts and row counts, among others
- Not rolled back as part of rollback tran
- sp\_configure “housekeeper free write percent” **must be at least 1**
  - Otherwise systabstats is only refreshed by
    - update % statistics
    - exec sp\_flushstats
  - FYI value of 1 is rarely high enough, consider 2-5

# When stats weren't updated by update stats

---

- Stats are only updated by the same command that generated them
  - If you run `update all statistics` once
  - Thereafter run `update index statistics` weekly
  - The stats on unindexed columns are **not** updated
  - But it's even worse than that...

# When stats weren't updated by update stats

---

- Some attributes of stats have “stickiness”
  - Once set explicitly, they are preserved unless set explicitly again
  - If you don't set them, future commands inherit the values

```
update index statistics tableA using 50 values
```

```
update index statistics tableA -- also 50 values
```

# When stats weren't updated by update stats

---

- Which stat attributes are sticky (by default) if explicitly set?
  - using NNN values (# of steps)
  - histogram tuning factor
  - sampling
  - hashing / partial\_hashing / no\_hashing

# When stats weren't updated by update stats

---

- How to change the value of sticky attributes?
  - A new command with explicitly different values
  - Remove the stats completely
  - Remove the stickiness

# When stats weren't updated by update stats

---

- Stats can be deleted explicitly with `delete stats`
  - This is good practice before every `update stats` command
  - Backup the stats first with `optdiag`
  - This only deletes column distribution stats from `sysstatistics`
    - Does not delete table-level stats from `systabstats`



# When stats weren't updated by update stats

---

- Stats can sometimes be deleted implicitly
  - Drop & recreate index
    - But only if the columns don't appear in any other index
    - Tip: use 0 steps to maintain stats when dropping & recreating
    - Useful when dropping & recreating as second-class reorg
    - ... but reorg rebuild online / defrag are both better

# When stats weren't updated by update stats

---

- Remove stickiness...
  - ... on this table
    - `sp_modifystats ... REMOVE_STICKINESS`
  - ... on every table
    - `sp_configure "enable sticky statistics", 0`
- This is all or nothing
  - Can't specify that some attributes are sticky and others are not

# Update stats for free

---

- Update stats for free if table already has data
- Create NC index (or DOL CL index) = update statistics ... index ...
- Create CL index (APL) = update statistics ... table
- What's missing?
  - Non-leading columns
    - We cannot do anything about this
  - Type or attributes of the stats
    - We do have some control over these

# Update stats for free

```
create index ... with statistics using NNN values
,          statistics hashing
,          no_hashing
,          new_hashing
,          statistics histogram_tuning_factor = N
,          statistics max_resource_granularity = N
,          statistics print_progress = N
```

- Only works when the table has data

# Update stats... in the right ways

---

- Which version
- Consumers
- Sampling
- Hashing
- Steps / histogram tuning factor
- Out of range values / dynamic histograms
- Under the hood: sorts and scans
- Avoiding duplicate work / minimum set coverage

# Which version: there is only update statistics

---

- No such thing as update index statistics, or update all statistics
  - Internally transformed into multiple update statistics commands
  - Implications for duplicate effort
  - Implications for minimum set coverage, more on this later

# Which version: there is only update statistics

- update index statistics tableA

Example:

tableA (a, b, c, d, e, f, g)

index1 (a, b, c)

index2 (a, d, e)

index3 (e)

# Which version: there is only update statistics

- update index statistics tableA
  - update index statistics tableA index1
  - update index statistics tableA index2
  - update index statistics tableA index3

Example:

tableA (a, b, c, d, e, f, g)

index1 (a, b, c)

index2 (a, d, e)

index3 (e)



# Which version: there is only update statistics

- update index statistics tableA
  - update index statistics tableA index1
    - update statistics tableA index1
    - update statistics tableA (b)
    - update statistics tableA (c)
  - update index statistics tableA index2
    - update statistics tableA index2
    - update statistics tableA (d)
    - update statistics tableA (e)
  - update index statistics tableA index3
    - update statistics tableA index3

-- a

-- b

-- c

-- a

-- d

-- e

-- e

Example:

tableA (a, b, c, d, e, f, g)

index1 (a, b, c)

index2 (a, d, e)

index3 (e)

# Which version: there is only update statistics

- update all statistics tableA

Example:

tableA (a, b, c, d, e, f, g)

index1 (a, b, c)

index2 (a, d, e)

index3 (e)

# Which version: there is only update statistics

- update all statistics tableA
  - update index statistics tableA index1
  - update index statistics tableA index2
  - update index statistics tableA index3
  - update statistics tableA (f)
  - update statistics tableA (g)

Example:

tableA (a, b, c, d, e, f, g)

index1 (a, b, c)

index2 (a, d, e)

index3 (e)

# Which version: there is only update statistics

- update all statistics tableA
  - update index statistics tableA index1
    - update statistics tableA index1 -- a
    - update statistics tableA (b) -- b
    - update statistics tableA (c) -- c
  - update index statistics tableA index2
    - update statistics tableA index2 -- a
    - update statistics tableA (d) -- d
    - update statistics tableA (e) -- e
  - update index statistics tableA index3
    - update statistics tableA index3 -- e
  - update statistics tableA (f) -- f
  - update statistics tableA (g) -- g

Example:

tableA (a, b, c, d, e, f, g)

index1 (a, b, c)

index2 (a, d, e)

index3 (e)

# Consumers

---

- In theory we can divide up the work and do it in parallel
- Many prerequisites for consumers =  $N$ 
  - Must have  $N+1$  “number of worker processes” available at runtime
  - Must have “max parallel degree”  $> N+1$
  - Must have “max utility parallel degree”  $> N+1$
- Incompatible with sampling or hashing

# Sampling

---

- All the pain of update statistics is based on number of pages
- Read fewer pages = less pain... but stats quality suffers...
- update statistics ... with sampling = 10 percent
- sp\_configure “sampling percent”, 10
- Incompatible with consumers or hashing

# Sampling

---

- The devil is in the details
  - Sampling **does not update density** from any non-sampling run
    - `update table statistics tableA`
  - Sampling only updates density if stats are deleted first
  - Sampling only applies to running on individual columns
    - Ignored for running on index or table
  - Sampling only applies to data pages
    - Ignored for all index page scans

# Sampling

---

- Implications for density as it is not adjusted for sampling
- Consider 1000 rows all with same value in this column = density 1.0
- Run update stats with sampling = 10 percent
- ASE will see 10% x 1000 rows = 100 rows all the same, out of 1000 rows
- ASE will mark density = 0.1 !!!
- If you still use sampling, consider sp\_modifystats
  - Adjust densities by  $\frac{1}{\text{sampling}}$
  - Express percentage here as a decimal 0.0 – 1.0



# Hashing

---

- More accurate than sampling, and probably less work
  - Far less tempdb disk space, far less procedure cache
  - A lot more data cache (in whichever case tempdb uses)
- Incompatible with consumers or sampling
- Consumers and sampling are dead, use hashing instead

# Hashing

---

- with hashing | partial\_hashing | no\_hashing
  - hashing = hash regardless of cardinality
  - partial\_hashing = hash up to 64K unique values, else sort as usual
    - If existing stats say > 64K unique values, no attempt to hash!
- sp\_configure “update statistics hashing”
  - Controls server-wide behaviour during update statistics
- sp\_configure “utility statistics hashing”
  - Controls whether “free” stats during create index use hashing

# Steps and histogram tuning factor

---

- Number of steps = **requested** number of histogram steps
- Histogram tuning factor = multiply # steps by this number **if it helps**
- Actual number of steps could be lower or higher than requested
  - 1 or 2 steps (NULL) if all values in column are duplicates
  - Up to *# of steps \* histogram tuning factor*
- Advice from SAP: 1000 total is good, 2000 total is too many
  - With default htf=20, implies request 50 steps

# Steps and histogram tuning factor

---

- Maybe better: 100 steps, htf = 10, set globally via sp\_configure
- Histogram tuning factor not guaranteed to increase steps
- Neither is requesting steps... but is a lot more likely to
- The average number of steps is higher for 100..1000 than for 20..1000
- Remember both attributes are sticky!

# When the histogram can't keep up with the data

---

- Two solutions for how to handle column values that only increase
- Default behaviour: out of range SARGs have selectivity = 0 (!!)
  - This can significantly alter query plans

# When the histogram can't keep up with the data

---

- ASE 15.0.2 ESD#2: out of range histograms
- Extrapolate OOR value based on existing histogram behaviour
- Enabled by default server-wide
- Disable server-wide with TF 15355 (documented & supported)
- Disable per column with
  - `update statistics table (column) out_of_range off`
  - Does not actually run update stats; toggles a status
  - May occasionally resolve weird optimiser behaviour

# When the histogram can't keep up with the data

---

- ASE 16.0 SP01: dynamic histograms
  - Documentation bug: not mentioned in some docs > 16.0 SP01
- Update the very last histogram cell transactionally
  - Hold in memory only
  - Flushed at polite shutdown or “sysstatistics flush interval”
    - Remember latter default is zero

# When the histogram can't keep up with the data

---

- ASE 16.0 SP01: dynamic histograms
- Enable per column with
  - `update statistics table (column) using dynamic histogram on`
  - Does not actually run update stats; toggles a status
- Might resolve poor optimiser choices when column only increases
- Definitely makes things worse if column doesn't only increase
  - There is overhead per insert / update / delete



# Under the hood... scans and locks

Command	Locking scheme	Index	Scans	Sorts (tempdb)	Locking
update statistics tbl (col)	APL	--	Table scan into worktable	Sorts the worktable	Sh Int Tbl, Sh Page
	DOL	--	Table scan into worktable	Sorts the worktable	Level 0 dirty reads
update statistics tbl index	APL	NC	Leaf level index scan	--	Sh Int Tbl, Sh Page
	APL	CL	Table scan	--	Sh Int Tbl, Sh Page
	DOL	any	Leaf level index scan	--	Level 0 dirty reads
update statistics tbl	APL	any	Table scan + leaf-level index scan per NC index	--	Sh Int Tbl, Sh Page
	DOL	any	Table scan + leaf-level index scan per index	--	Level 0 dirty reads

- Specifying consumers always applies to sorts on worktables
- Specifying consumers only applies to base table if it is partitioned

# Under the hood... hashing

Command	Cardinality	Sampling	Partial_hashing	Hashing
update statistics tbl (col)	low	Data only	One hash on column	One hash on column
	high	Data only	One worktable + sort on column	One hash on column
update statistics tbl index	low	Ignored	One hash on index	One hash on index
	high	Ignored	One worktable + sort on index leading column	One hash on index
update statistics tbl	low	Ignored	One hash per index	One hash per index
	high	Ignored	One worktable + sort on index leading column, per index	One hash per index
update index statistics	low	Ignored	One hash per index	One hash per index
	high	Ignored	One worktable + sort per column	One hash per index
update all statistics	low	Ignored	One hash (!)	One hash (!!)
	high	Ignored	One worktable + sort per column	One hash (!!!)

# Avoiding duplicate work

---

- Remember everything is broken down into multiple update statistics
- This means a duplicate scan every time a column appears
- ... unless you use (full) hashing!
- `update all statistics ... with hashing` uses a single scan!
  - Also little tempdb space and procedure cache
  - But a lot of tempdb data cache
- What to do if this isn't possible or practical?

# Avoiding duplicate work

- Minimal set coverage is a famously hard problem
- update statistics = stats on a, e  
= 2 scans
- update index statistics = stats on a, b, c, d, e  
= 7 scans
- update all statistics = stats on a, b, c, d, e, f, g  
= 9 scans
- update statistics (a, b, c, d, e) = stats on a, b, c, d, e  
= 5 scans

Example:

tableA (a, b, c, d, e, f, g)

index1 (a, b, c)

index2 (a, d, e)

index3 (e)

# So what is “the right way”?

---

- DON'T ever run update stats in PROD; optdiag stats from a copy
- DO use datachange() while understanding its limits
- DO use optdiag out to backup stats before any delete stats
- DO use delete stats before any update stats
- DO understand stat stickiness; recommend disabling
- DON'T use consumers
- DON'T use sampling
- DO use sp\_modifystats on densities if you still use sampling

# So what is “the right way”?

---

- DO set “sysstatistics flush interval” to 15 or 60
- DO set a sensible global value for number of steps
- DO set a sensible global value for histogram tuning factor
- DO use hashing wherever possible
- DO use minimal set coverage where full hashing is not possible
  - Superior to update index statistics !

# Q & A, and thank you

---

Joe Woodhouse

joe.woodhouse@primadonnaconsulting.com

Too busy putting out fires to reduce your toil?

Answering the on-call phone too often?

Drowning in technical debt?

Joe is a freelance consultant available through Prima Donna Consulting and can be engaged flexibly.

Not a lawyer – no charge for emails!

© Prima Donna Consulting 2023

These materials are provided for informational purposes only, without representation or warranty of any kind, and Prima Donna Consulting shall not be liable for errors or omissions.